

FIT2086 Assignment 3

Due Date: 11:59PM, Sunday, 8/11/2019

Introduction

There are total of three questions worth $10 + 18 + 14 = 42$ marks in this assignment.

This assignment is worth a total of 20% of your final mark, subject to hurdles and any other matters (e.g., late penalties, special consideration, etc.) as specified in the FIT2086 Unit Guide or elsewhere in the FIT2086 Moodle site (including Faculty of I.T. and Monash University policies).

Students are reminded of the Academic Integrity Awareness Training Tutorial Activity and, in particular, of Monash University's policies on academic integrity. In submitting this assignment, you acknowledge your awareness of Monash University's policies on academic integrity and that work is done and submitted in accordance with these policies.

Submission: No files are to be submitted via e-mail. Correct files are to be submitted to Moodle, as given above. You must submit the following three files:

1. One PDF file containing non-code answers to all the questions that require written answers. This file should also include all your plots.
2. The two required R script files containing R code answers as discussed in Question 2 and 3.

Please read these submission instructions carefully and take care to submit the correct files in the correct places.

Question 1 (10 marks)

This question will require you to analyse a regression dataset. The file `housing.ass3.2020.csv` contains the data that we will use for this question. This dataset is a modified version of the Boston housing data which was collected to study house prices in the metropolitan region of Boston. In this data set, each observation represents a particular suburb from the Boston region. The outcome, `medv`, is the median value of owner-occupied homes in 1,000 in the suburb. The variables are summarised in Table 1. The data consists of $p = 12$ variables measured on $n = 250$ suburbs. We are interested in discovering which predictors are good determinants of housing price, and how these variables effect the median house price.

1. Fit a multiple linear model to the housing data using R. Using the results of fitting the linear model, which predictors do you think are possibly associated with median house value, and why? Which three variables appear to be the strongest predictors of housing price, and why? **[3 marks]**
2. How would your assessment of which predictors are associated change if you used the Bonferroni procedure with $\alpha = 0.05$? **[1 mark]**
3. Describe what effect the per-capita crime rate (`crim`) appears to have on the median house price. Describe what effect a suburb having frontage on the Charles River has on the median house price for that suburb. **[2 marks]**
4. Use the stepwise selection procedure, with the BIC criterion, to prune out potentially unimportant variables. Write down the final regression equation obtained after pruning. **[1 mark]**
5. If a council wanted to try and improve the median house value in their suburb, what does the model that we found in Question 1.4 suggest they could try and do? **[2 marks]**
6. Table 2 gives the values of predictors for a new suburb. Use the model found in Question 1.4 to predict the median house price for this suburb. Provide a 95% confidence interval for this prediction. **[1 mark]**

Variable name	Description	Values
crim	Per-capita crime rate	> 0
zn	Proportion of residential land zoned for lots over 25,000 sq. ft.	$0 - 100$
indus	Proportion of non-retail business acres per town	$0 - 100$
chas	Does the suburb front the Charles River?	$0 = \text{No}, 1 = \text{Yes}$
nox	Nitric oxides concentration (parts per 10 million)	> 0
rm	Average number of rooms per dwelling	≥ 1
age	Proportion of owner-occupied units built prior to 1940	$0 - 100$
dis	Weighted distances to five Boston employment centres	> 0
rad	Index of accessibility to radial highways	> 0
tax	Full-value property-tax rate per \$10,000	$187 - 711$
ptratio	Pupil-teacher ratio	> 0
lstat	Percentage of “lower status” of the population	$0 - 100$
medv	Median value of owner-occupied homes in \$1,000s	> 0

Table 1: Boston Housing Data Dictionary.

Variable	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	lstat
Value	0.04741	0	11.93	0	0.573	6.03	80.8	2.505	1	273	21	7.88

Table 2: Boston Housing Data Dictionary.

Question 2 (18 marks)

In this question we will analyse the data in `heart.train.ass3.2020.csv`. In this dataset, each observation represents a patient at a hospital that reported showing signs of possible heart disease. The outcome is presence of heart disease (HD), or not, so this is a classification problem. The predictors are summarised in Table 3. We are interested in learning a model that can predict heart disease from these measurements. To answer this question you must:

- Provide an R script containing all the code you used to answer the questions. Please use comments to ensure that the code used to identify each question is **clearly identifiable**. Call this `fn.sn.Q2.R`, where “fn.sn” is your first name followed by your family name.
- Provide appropriate written answers to the questions, along with any graphs, in a report document.

When answering this question, you must use the `rpart` package that we used in Studio 9. The wrapper function for learning a tree using cross-validation that we used in Studio 9 is contained in the file `wrappers.R`. Don't forget to source this file to get access to the function.

1. Using the techniques you learned in Studio 9, fit a decision tree to the data using the `tree` package. Use cross-validation with 10 folds and 5,000 repetitions to select an appropriate size tree. What variables have been used in the best tree? How many leaves (terminal nodes) does the best tree have? [2 marks]
2. Plot the tree found by CV, and discuss clearly and thoroughly in plain English what it tells you about the relationship between the predictors and heart disease. (*hint: you can use the `text(cv$best.tree,pretty=12)` function to add appropriate labels to the tree*). [3 marks]
3. For classification problems, the `rpart` package only labels the leaves with the most likely class. However, if you examine the tree structure in its textual representation on the console, you can determine the probabilities of having heart disease (see Question 2.3 from Studio 9 as a guide) in each leaf (terminal node). Take a screen-capture of the plot of the tree (don't forget to use the “zoom” button to get a larger image) or save it as an image using the “Export” button in R Studio.

Then, use the information from the textual representation of the tree available at the console and annotate the tree in your favourite image editing software; next to all the leaves in the tree, add text giving the probability of contracting heart disease. Include this annotated image in your report file. [2 marks]
4. According to your tree, which predictor combination results in the highest probability of having heart-disease? [1 mark]
5. We will also fit a logistic regression model to the data. Use the `glm()` function to fit a logistic regression model to the heart data, and use stepwise selection with the BIC score to prune the model. What variables does the final model include, and how do they compare with the variables used by the tree estimated by CV? Which predictor is the most important in the logistic regression? [3 marks]
6. Write down the regression equation for the logistic regression model you found using step-wise selection. [1 mark]

7. The file `heart.test.ass3.2020.csv` contains the data on a further $n' = 200$ individuals. Using the `my.pred.stats()` function contained in the file `my.prediction.stats.R`, compute the prediction statistics for both the tree and the step-wise logistic regression model on this test data. Contrast and compare the two models in terms of the various prediction statistics? Would one potentially be preferable to the other as a diagnostic test? Justify your answer. [2 marks]
8. Calculate the *odds* of having heart disease for the patient in the 69th row of the test dataset. The odds should be calculated for both:
- (a) the tree model found using cross-validation; and
 - (b) the step-wise logistic regression model.

How do the predicted odds for the two models compare? [2 marks]

9. For the logistic regression model using the predictors selected by BIC in Question 2.6, use the bootstrap procedure (use at least 5,000 bootstrap replications) to find a confidence interval for the probability of having heart disease for patient in the 69th row in the test data. Use the `bca` option when computing this confidence interval. Discuss this confidence interval in comparison to the predicted probabilities of having heart disease for both the logistic regression model and the tree model. [2 marks]

Variable name	Description	Values
AGE	Age of patient in years	29 – 77
SEX	Sex of patient	M = Male F = Female
CP	Chest pain type	Typical = Typical angina Atypical = Atypical angina NonAnginal = Non anginal pain Asymptomatic = Asymptomatic pain
TRESTBPS	Resting blood pressure (in <i>mmHg</i>)	94 – 200
CHOL	Serum cholesterol in <i>mg/dl</i>	126 – 564
FBS	Fasting blood sugar > 120 <i>mg/dl</i> ?	<120 = No >120 = Yes
RESTECG	Resting electrocardiographic results	Normal = Normal ST.T.Wave = ST wave abnormality Hypertrophy = showing probable hypertrophy
THALACH	Maximum heart rate achieved	71 – 202
EXANG	Exercise induced angina?	N = No Y = Yes
OLDPEAK	Exercise induced ST depression relative to rest	0 – 6.2
SLOPE	Slope of the peak exercise ST segment	Up = Up-sloping Flat = Flat Down = Down-sloping
CA	Number of major vessels colored by flourosopy	0 – 3
THAL	Thallium scanning results	Normal = Normal Fixed.Defect = Fixed fluid transfer defect Reversible.Defect = Reversible fluid transfer defect
HD	Presence of heart disease	N = No Y = Yes

Table 3: Heart Disease Data Dictionary. ST depression refers to a particular type of feature in an electrocardiograph (ECG) signal during periods of exercise. Thallium scanning refers to the use of radioactive Thallium to check the fluid transfer capability of the heart.

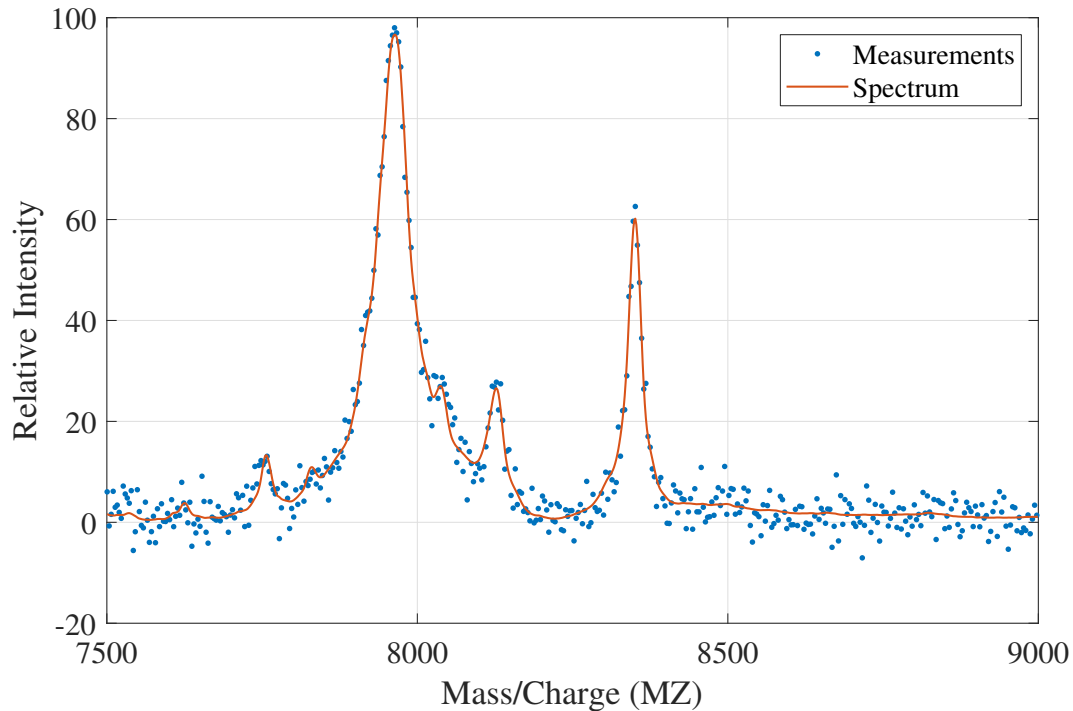


Figure 1: Noisy measurements from a (simulated) mass spectrometry reading. The “true” (unknown) measurements are shown in orange, and the noisy measurements are shown in blue.

Question 3 (14 marks)

Data Smoothing

Data “smoothing” is a very common problem in data science and statistics. We are often interested in examining the unknown relationship between a dependent variable (y) and an independent variable (x), under the assumption that the dependent variable has been imperfectly measured and has been contaminated by measurement noise. The model of reality that we use is

$$y = f(x) + \varepsilon$$

where $f(x)$ is some unknown, “true”, potentially non-linear function of x , and $\varepsilon \sim N(0, \sigma^2)$ is a random disturbance or error. This is called the problem of function estimation, and the process of estimating $f(x)$ from the noisy measurements y is sometimes called “smoothing the data” (even if the resulting curve is not “smooth” in a traditional sense, it is less rough than the original data).

In this question you will use the k -nearest neighbours machine learning technique to smooth data. This technique is used frequently in practice (think for example the 14-day rolling averages used to estimate coronavirus infection numbers). This question will explore its effectiveness as a smoothing tool.

Mass Spectrometry Data Smoothing

The file `ms.train.2020.csv` contains $n = 443$ measurements from a mass spectrometer. Mass spectrometry is a chemical analysis tool that provides a measure of the physical composition of a material. The outputs of a mass spectrometry reading are the intensities of various ions, indexed by their mass-to-charge ratio. The resulting spectrum usually consists of a number of relatively sharp peaks that indicate a concentration of particular ions, along with an overall background level. A standard problem is that the measurement process is generally affected by noise – that is, the sensor readings are imprecise and corrupted by measurement noise. Therefore, smoothing, or removing the noise is crucial as it allows us to get a more accurate idea of the true spectrum, as well as determine the relative quantity of the ions more accurately. However, we would *ideally* like for our smoothing procedure to not damage the important information contained in the spectrum (i.e., the heights of the peaks).

The file `ms.train.csv` contains measurements of our mass spectrometry reading; `ms.train$MZ` are the mass-to-charge ratios of various ions, and `ms.train$intensity` are the measured (noisy) intensities of these ions in our material. The file `ms.test.2020.csv` contains $n = 886$ different values of MZ along with the “true” intensity values, stored in `ms.test.2020$intensity`. These true values have been found by using several advanced statistical techniques to smooth the data, and are being used here to see how close your estimated spectrum is to the truth. For reference, the samples `ms.train$intensity` and the value of the true spectrum `ms.test$intensity` are plotted in Figure 1 against their respective MZ values. To answer this question you must:

- Provide an R script containing all the code you used to answer the questions. Please use comments to ensure that the code used to identify each question is **clearly identifiable**. Call this file `fn.sn.Q3.R`, where “`fn.sn`” is your first name followed by your family name.
- Provide appropriate written answers to the questions, along with any graphs, in a non-handwritten report document.

To answer this question, you must use the `kkn` and `boot` packages that we used in Studios 9 and 10.

Questions

1. Use the k -nearest neighbours method (k -NN) to estimate the underlying spectrum from the training data. Use the `kkn` package we examined in Studio 9 to provide predictions for the MZ values in `ms.test`, using `ms.train` as the training data. You should use the `kernel = "optimal"` option when calling the `kkn()` function. This means that the predictions are formed by a weighted average of the k points nearest to the point we are trying to predict, the weights being determined by how far away the neighbours are from the point we are trying to predict.
 - (a) For each value of $k = 1, \dots, 25$, use k -NN to estimate the values of the spectrum for the MZ values in `ms.test$MZ`. Then, compute the mean-squared error between your estimates of the spectrum, and the true values in `ms.test$intensity`. Produce a plot of these errors against the various values of k . [1 mark]
 - (b) Produce four graphs, each one showing: (i) the training data points (`ms.train$intensity`), (ii) the true spectrum (`ms.test$intensity`) and (iii) the estimated spectrum (predicted `intensity` values for the MZ values in `ms.test.csv`) produced by the k -NN method for four different values of k ; do this for $k = 2$, $k = 5$, $k = 10$ and $k = 25$. Make sure the graphs have clearly labelled axis' and a clear legend. Use a different colour for your estimated curve. [3 marks]
 - (c) Discuss, qualitatively, and quantitatively (in terms of mean-squared error on the true spectrum) the four different estimates of the spectrum. [2 marks]

2. Use the cross-validation functionality in the `kknn` package to select an estimate of the best value of k (make sure you still use the `optimal` kernel). What value of k does the method select? How does it compare to the (in practice, unknown) value of k that would minimise the actual mean-squared error (as computed in Question 3.1a)? [1 mark]
3. Using the estimates of the curve produced in the previous question, see if you can provide an estimate of the variance of the sensor/measurement noise that has corrupted our intensity measurements. [1 mark]
4. Do any of the estimated spectra achieve our aim of providing a smooth, low-noise estimate of background level as well as accurate estimation of the peaks? Explain why you think the k -NN method is able to achieve, or not achieve, this aim. [2 marks]
5. An important task when processing mass spectrometry signals is to locate the peaks, as this gives information on which elements are present. From the smoothed signal produced using the value of k found in Question 3.2, which value of MZ corresponds to the maximum estimated abundance? [1 mark]
6. Using the bootstrap procedure (use at least 5,000 bootstrap replications), write code to find a confidence interval for the k -nearest neighbours estimate of relative abundance at a specific MZ value. Use this code to obtain a 95% confidence interval for the estimate of relative abundance at the MZ value you determined previously in Question 3.5 (i.e., the value corresponding to the highest relative intensity). Compute confidence intervals using the k determined in Question 3.2, as well as $k = 3$ neighbour and $k = 20$ neighbours. Report these confidence intervals. Explain why you think these confidence intervals vary in size for different values of k . [3 marks]